
Exercise Sheet Deep Learning

Part 3: Attacks and Defenses

Summer 22

This sheet includes a theoretical part and a practical assignment on the third part of the lecture Deep Learning (3_Adversarials). Both parts give 20 points maximum each. Please hand in solutions as a pdf in groups of at most three persons via LernraumPlus.

name1:

name2:

name3:

PART I – THEORY: For the following, you might answer only YES/NO (or abstain), or you can add short arguments (at most two lines per question). If you are not sure, it is better to abstain.

1. For the following attacks, it holds:

☐ **yes** ☐ **no** Backdoor attacks need access to the training scheme. They try to change the function such that the network can no longer be used.

☐ **yes** ☐ **no** Data poisoning attacks try to change the data such that special triggers are included.

☐ **yes** ☐ **no** Model inversion attacks can be countered using mechanisms of differential privacy

☐ **yes** ☐ **no** Model inversion attacks need to have access to the model function.

2. For the following adversarial attacks, it holds:

☐ **yes** ☐ **no** Fast gradient sign and basic iterative method are based on gradient schemes.

☐ **yes** ☐ **no** Deepfool uses an adaptive step size.

☐ **yes** ☐ **no** Universal adversarial perturbations incorporate a regularization against translations of visual objects in scenes.

☐ **yes** ☐ **no** Adversarial attacks need to have access to the model gradient or estimate it implicitly

3. Defenses against attacks ...

- ☐yes ☐no Data poisoning can be detected by clustering the training set.
- ☐yes ☐no Homomorphic privacy enables the public release of models trained on personal data.
- ☐yes ☐no Adversarial training solves the adversarial training loss exactly via computing worst case adversarials in the inner loop.
- ☐yes ☐no Certified defenses can rely on Lipschitz bounds.

4. Attacks in reality:

- ☐yes ☐no Some training data might deteriorate model accuracy rather than helping it
- ☐yes ☐no Attacks need to address the whole input space
- ☐yes ☐no Attacks can be designed such that they hold for different object view points or different scenes in a stream
- ☐yes ☐no Data compression can provably avoid attacks.

5. The following holds:

- ☐yes ☐no Robustified networks have the same accuracy as original ones
- ☐yes ☐no Adversarial risk and minimum perturbation risk are identical for the mean squared error.
- ☐yes ☐no Adversarial training complexity strongly depends on the design of adversarial attacks in the inner loop
- ☐yes ☐no Adversarial examples do not exist for linear models.

PARTII – PRACTICE: You can use code and models which are publicly available, please clearly reference such sources. It might be a good idea to start with the examples given in the practical part of the lecture (available at <https://jgoepfert.pages.ub.uni-bielefeld.de/talk-deep-learning>) or you might want to resort to popular toolboxes such as <https://adversarial-robustness-toolbox.readthedocs.io/en/latest/>. Please give a link to your code, and please describe the experiments and results of your approach in a pdf which is well structured (e.g. modeling/training parameters/training/results/interpretation, use itemize, keywords are fine) and enables reproducibility as well as easy access to your main results. Please use at most one page for both practical parts together including graphs and images.

1. Take the model for the FashionMNIST data from the first sheet. Perform targeted and untargeted attacks for different sampled examples and classes and report, whether some attacks are harder to achieve than others (depending e.g. in classes, location of data, ...) and why – you might want to visualize the classes using tSNE or UMAP for this purpose or even use DeepView.
2. Vary the attacks (i) such that the misclassification of the attack is done with a large confidence (confidence score coming directly from the soft-max for simplicity), and (ii) such that the attack pattern works for more than one data point, i.e. it attacks several input signals at once.